

```
In [1]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import (
    accuracy_score,
    classification_report,
    confusion_matrix,
    ConfusionMatrixDisplay
)
import matplotlib.pyplot as plt
```

```
In [3]: df = pd.read_csv(r"C:\Users\joshb\Downloads\imdb_reviews_sample.csv")

# Convert text labels ("pos"/"neg") to numeric labels (1/0)
df["label"] = df["sentiment"].map({"pos": 1, "neg": 0})
```

```
In [5]: #split the data
X_train, X_test, y_train, y_test = train_test_split(
    df[["review"]], df[["label"]], test_size=0.3, random_state=42
)
```

```
In [7]: #Convert reviews into TF-IDF feature vectors

vectorizer = TfidfVectorizer(stop_words="english", max_features=5000)
X_train_tfidf = vectorizer.fit_transform(X_train)
X_test_tfidf = vectorizer.transform(X_test)
```

```
In [9]: log_reg = LogisticRegression(max_iter=200)
log_reg.fit(X_train_tfidf, y_train)
y_pred_log = log_reg.predict(X_test_tfidf)

print(" Logistic Regression Accuracy:", round(accuracy_score(y_test, y_pred_log), 3))
print(classification_report(y_test, y_pred_log))
```

```
Logistic Regression Accuracy: 0.828
      precision    recall  f1-score   support
          0       0.85     0.81     0.83     305
          1       0.81     0.85     0.83     295

      accuracy                           0.83     600
     macro avg       0.83     0.83     0.83     600
weighted avg       0.83     0.83     0.83     600
```

```
In [11]: #Train Naïve Bayes model

nb = MultinomialNB()
nb.fit(X_train_tfidf, y_train)
y_pred_nb = nb.predict(X_test_tfidf)


print("\n Naïve Bayes Accuracy:", round(accuracy_score(y_test, y_pred_nb), 3))
print(classification_report(y_test, y_pred_nb))
```

```
Naïve Bayes Accuracy: 0.835
      precision    recall  f1-score   support
          0       0.82     0.87     0.84     305
          1       0.85     0.80     0.83     295

      accuracy                           0.83     600
     macro avg       0.84     0.83     0.83     600
weighted avg       0.84     0.83     0.83     600
```

```
In [13]: # Show Confusion Matrix (Logistic Regression)
cm = confusion_matrix(y_test, y_pred_log)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Negative", "Positive"])
disp.plot(cmap="Blues")
plt.title("Confusion Matrix - Logistic Regression")
plt.show()
```

Confusion Matrix - Logistic Regression



In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js